

VASISHTA CHANDALA

Data Scientist | Applied AI & Data Engineering

vasishta.chandala@gmail.com

www.linkedin.com/in/vasishta-chandala | github.com/vasishta02 | Portfolio

SUMMARY

Data Scientist and startup founding-team contributor specializing in predictive modeling, growth analytics, and cloud data engineering. Builds Python/SQL pipelines, deploys ML workflows, and evaluates grounded AI applications to improve operational quality, subscriber retention, and efficiency.

EXPERIENCE

H.A.D.E.S | Data Scientist, Founding Team Member | Jan 2024 - Feb 2026

Remote | Core role through Jul 2025; ad hoc support through Feb 2026.

- Built **AWS Kinesis and Lambda** ingestion processing **10K+ events/minute** for operational monitoring, forecasting, and downstream analytics.
- Reduced data latency by **40%** through data preparation and analytics workflows spanning Glue, S3, PostgreSQL, dbt, Airflow, and Snowflake.
- Developed **XGBoost defect prediction across 50+ variables**, demand forecasting, and anomaly detection; deployed ML workflows with Docker and SageMaker, contributing to **20% higher quality/output** and **15% fewer recurring issues**.
- Translated founder and CEO priorities into **50+ KPIs** and Tableau/Power BI dashboards; analytics and automation contributed to **\$50K+ annual savings**.

Hive Media | Data Scientist, Growth Analytics | Jul 2023 - Dec 2023

- Analyzed subscription behavior across **500K+ users and 1TB+ monthly data**; built churn/LTV analyses, cohorts, and segmentation that informed initiatives achieving **25% higher retention** and **18% lower churn**.
- Designed and analyzed **A/B tests** for product and content changes, supporting **15-20% higher engagement**; monetization and relevance analysis contributed to **18% higher ad relevance**.
- Built recurring behavioral analytics using Python, SQL, Spark, Hadoop, and Hive/HiveQL to support growth and monetization decisions.

COEMPT EDUTECK | Data Analyst Intern | Jan 2023 - Jul 2023

- Validated **100K+ records** with Python/SQL and automated KPI reporting, improving efficiency by **22%** and saving **~30 hours/month**; performed data-quality checks and trend analysis.

SELECTED PROJECTS

ClimateGPT Fusion | Hybrid RAG + MCP Climate Analytics | 2025

- Built hybrid RAG + MCP analytics over **EDGAR v2024 (2000-2024)**, combining entity normalization, parameterized SELECT-only SQL, schema validation, and Streamlit/DuckDB to ground answers in traceable data.
- Implemented a **50-question evaluation harness**: hallucinations fell **9% to 0%**, query success rose **94% to 100%**, and provenance completeness improved **35 to 92/100**; average response time was **~5.7 seconds**.

Bridge Condition Prediction (Bridges at Risk) | 2024

- Prepared **600K+ infrastructure records** with PySpark/Databricks and Snowflake; applied XGBoost and SHAP for interpretable prediction, reporting **77.7% accuracy** and four clusters with a **0.87 silhouette score**.

GridSense | Multimodal ML for Grid Asset Risk | 2026

- Built **PyTorch/ResNet18** image, text, and fusion models with SHAP; packaged FastAPI inference in Docker, added GitHub Actions/Kubernetes manifests, and implemented local LangChain/FAISS retrieval.

TECHNICAL SKILLS

Analytics & ML: Python, SQL, R, scikit-learn, XGBoost, PyTorch, SHAP, forecasting, A/B testing.

AI & Data Engineering: RAG, MCP, Spark/PySpark, Airflow, dbt, Snowflake, PostgreSQL, DuckDB.

Cloud & Delivery: AWS, Azure, BigQuery, Docker, GitHub Actions, FastAPI, Tableau, Power BI.

EDUCATION & CERTIFICATIONS

George Mason University - M.S. Data Analytics Engineering, Dec 2025 | GPA: 3.97/4.0

KL University - B.Tech Computer Science (Data Analytics), May 2023

Microsoft: Azure Fundamentals (AZ-900), Azure AI Fundamentals (AI-900), Azure Developer (AZ-204), DevOps Engineer (AZ-400).