

VASISHTA CHANDALA

Data Scientist | Applied AI & Data Engineering

vasishta02@gmail.com | www.linkedin.com/in/vasishta-chandala

github.com/vasishta02 | vasishta02.github.io/Portfolio/

SUMMARY

Data scientist building grounded AI applications, interpretable machine learning models, and cloud data pipelines. Experience across startup operations, subscription media, and analytics automation, connecting Python, SQL, and model evaluation to measurable business outcomes.

EXPERIENCE

H.A.D.E.S | Data Scientist, Founding Team Member | Jan 2024 - Feb 2026

Remote | Core role through Jul 2025; ad hoc support through Feb 2026.

- Built AWS Kinesis and Lambda workflows handling **10K+ events/minute**; developed Glue, S3, PostgreSQL, dbt, Airflow, and Snowflake analytics workflows that helped reduce data latency by **40%**.
- Developed XGBoost defect prediction using **50+ variables**, forecasting, and anomaly detection, contributing to **20% higher quality/output** and **15% fewer recurring issues**; packaged ML workflows with Docker and SageMaker.
- Delivered Tableau and Power BI dashboards tracking **50+ KPIs**, contributing to **\$50K+ annual savings**; collaborated with founders on KPI definitions and OpenAI API / TensorFlow prototypes.

Hive Media | Data Scientist, Growth Analytics | Jul 2023 - Dec 2023

- Analyzed **500K+ users and 1TB+ monthly data** using Python, SQL, and Spark; developed churn, LTV, cohort, and segmentation analyses supporting **25% higher retention** and **18% lower churn**.
- Designed and analyzed A/B tests supporting **15-20% higher engagement**; delivered monetization and content-relevance analysis, with Hadoop, HDFS, MapReduce, and Hive/HiveQL workflows.

COEMPT EDUTECK | Data Analyst Intern | Jan 2023 - Jul 2023

- Validated and analyzed **100K+ records** with Python and SQL; automated recurring reports, improving efficiency by **22%** and saving approximately **30 hours/month**.

SELECTED PROJECTS

ClimateGPT Fusion | Hybrid RAG + MCP Climate Analytics | 2025

- Integrated MCP, Streamlit, and DuckDB over EDGAR v2024 with entity normalization, SELECT-only SQL, and schema validation. A 50-question evaluation reported hallucinations **9% to 0%**, query success **94% to 100%**, and provenance completeness **35 to 92/100**.

Bridge Condition Prediction (Bridges at Risk) | 2024

- Combined PySpark, Databricks, Snowflake, XGBoost, and SHAP for infrastructure analytics; project results included **77.7% classification accuracy** and four clusters with a **0.87 silhouette score**.

GridSense | Multimodal ML for Grid Asset Risk | 2026

- Developed image, text, and fusion workflows using PyTorch / ResNet18 and SHAP; added FastAPI, Docker, GitHub Actions, Kubernetes manifests, Vertex AI job-spec scaffolding, and local LangChain / FAISS retrieval.

TECHNICAL SKILLS

Programming & ML: Python, SQL, R, scikit-learn, XGBoost, PyTorch, TensorFlow, SHAP, A/B testing.

AI & Data: RAG, MCP, DuckDB, Spark / PySpark, Airflow, dbt, Snowflake, PostgreSQL, BigQuery.

Cloud & Delivery: AWS, Azure, Docker, GitHub Actions, FastAPI, Streamlit, Tableau, Power BI.

EDUCATION & CERTIFICATIONS

George Mason University - M.S. Data Analytics Engineering, Dec 2025 | GPA: 3.97/4.0

KL University - B.Tech Computer Science (Data Analytics), May 2023

Microsoft: Azure Fundamentals (AZ-900), Azure AI Fundamentals (AI-900), Azure Developer (AZ-204), DevOps Engineer (AZ-400).